

# Importing necessary libraries

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import os

from keras.preprocessing.image import load_img, img_to_array
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Dense, Input, Dropout, GlobalAveragePooling2D, Flatten, Conv2D,
from keras.models import Model, Sequential
from keras.optimizers import Adam, SGD, RMSprop
```

```
2024-02-25 04:20:22.226513: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
2024-02-25 04:20:22.226607: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
2024-02-25 04:20:22.361731: E external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1515] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
```

```
In [2]: picture_size = 48
folder_path = "../input/face-expression-recognition-dataset/images/"
```

```
In [3]: expression = 'angry'

plt.figure(figsize=(12,12))
for i in range(1, 10, 1):
    plt.subplot(3,3,i)
    img = load_img(folder_path+"train/"+expression+"/"+
                    os.listdir(folder_path + "train/" + expression)[i], target_size=(
    plt.imshow(img)
plt.show()
```



Found 28821 images belonging to 7 classes.

Found 7066 images belonging to 7 classes.

## CNN Layers

```
In [5]: from keras.optimizers import Adam, SGD, RMSprop

no_of_classes = 7

model = Sequential()

#1st CNN Layer
model.add(Conv2D(64,(3,3),padding = 'same',input_shape = (48,48,1)))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

#2nd CNN Layer
model.add(Conv2D(128,(5,5),padding = 'same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout (0.25))

#3rd CNN Layer
model.add(Conv2D(512,(3,3),padding = 'same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout (0.25))

#4th CNN Layer
model.add(Conv2D(512,(3,3), padding='same'))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

model.add(Flatten())

#Fully connected 1st Layer
model.add(Dense(256))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))

# Fully connected Layer 2nd Layer
model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.25))
```

```
model.add(Dense(no_of_classes, activation='softmax'))
```

```
opt = Adam(lr = 0.0001)
```

```
model.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 48, 48, 64)	640
batch_normalization (Batch Normalization)	(None, 48, 48, 64)	256
activation (Activation)	(None, 48, 48, 64)	0
max_pooling2d (MaxPooling2D)	(None, 24, 24, 64)	0
dropout (Dropout)	(None, 24, 24, 64)	0
conv2d_1 (Conv2D)	(None, 24, 24, 128)	204928
batch_normalization_1 (Batch Normalization)	(None, 24, 24, 128)	512
activation_1 (Activation)	(None, 24, 24, 128)	0
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 128)	0
dropout_1 (Dropout)	(None, 12, 12, 128)	0
conv2d_2 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_2 (Batch Normalization)	(None, 12, 12, 512)	2048
activation_2 (Activation)	(None, 12, 12, 512)	0
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 512)	0
dropout_2 (Dropout)	(None, 6, 6, 512)	0
conv2d_3 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_3 (Batch Normalization)	(None, 6, 6, 512)	2048
activation_3 (Activation)	(None, 6, 6, 512)	0
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 512)	0
dropout_3 (Dropout)	(None, 3, 3, 512)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 256)	1179904

batch_normalization_4 (Batch Normalization)	(None, 256)	1024
activation_4 (Activation)	(None, 256)	0
dropout_4 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 512)	131584
batch_normalization_5 (Batch Normalization)	(None, 512)	2048
activation_5 (Activation)	(None, 512)	0
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 7)	3591

```

=====
Total params: 4478727 (17.08 MB)
Trainable params: 4474759 (17.07 MB)
Non-trainable params: 3968 (15.50 KB)

```

---

```

In [6]: from keras.optimizers import RMSprop, SGD, Adam
        from keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau

        checkpoint = ModelCheckpoint("./model.h5", monitor='val_acc', verbose=1, save_best_only=True,
        early_stopping = EarlyStopping(monitor='val_loss',
                                        min_delta=0,
                                        patience=3,
                                        verbose=1,
                                        restore_best_weights=True
                                        )

        reduce_learningrate = ReduceLROnPlateau(monitor='val_loss',
                                                factor=0.2,
                                                patience=3,
                                                verbose=1,
                                                min_delta=0.0001)

        callbacks_list = [early_stopping, checkpoint, reduce_learningrate]

        epochs = 48

        model.compile(loss='categorical_crossentropy',
                      optimizer = Adam(lr=0.001),
                      metrics=['accuracy'])

```

```

In [7]: history = model.fit_generator(generator=train_set,
                                      steps_per_epoch=train_set.n//train_set.batch_size,
                                      epochs=epochs,
                                      validation_data = test_set,
                                      validation_steps = test_set.n//test_set.batch_size,

```

```
callbacks=callbacks_list
)
```

```
/tmp/ipykernel_27/3853974908.py:1: UserWarning: `Model.fit_generator` is deprecated
and will be removed in a future version. Please use `Model.fit`, which supports gene
rators.
```

```
history = model.fit_generator(generator=train_set,
Epoch 1/48
```

```
2024-02-25 04:21:08.099823: E tensorflow/core/grappler/optimizers/meta_optimizer.cc:
961] layout failed: INVALID_ARGUMENT: Size of values 0 does not match size of permut
ation 4 @ fanin shape insequential/dropout/dropout/SelectV2-2-TransposeNHWCToNCHW-La
youtOptimizer
```

```
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
```

```
I0000 00:00:1708834872.928062      85 device_compiler.h:186] Compiled cluster using
XLA! This line is logged at most once for the lifetime of the process.
```

```
225/225 [=====] - 233s 995ms/step - loss: 1.7677 - accurac
y: 0.3225 - val_loss: 1.9443 - val_accuracy: 0.3612 - lr: 0.0010
```

```
Epoch 2/48
```

```
225/225 [=====] - 29s 129ms/step - loss: 1.4219 - accuracy:
0.4521 - val_loss: 1.5272 - val_accuracy: 0.4477 - lr: 0.0010
```

```
Epoch 3/48
```

```
225/225 [=====] - 30s 134ms/step - loss: 1.2736 - accuracy:
0.5118 - val_loss: 1.5533 - val_accuracy: 0.4402 - lr: 0.0010
```

```
Epoch 4/48
```

```
225/225 [=====] - 29s 129ms/step - loss: 1.1910 - accuracy:
0.5443 - val_loss: 1.1415 - val_accuracy: 0.5729 - lr: 0.0010
```

```
Epoch 5/48
```

```
225/225 [=====] - 30s 132ms/step - loss: 1.1326 - accuracy:
0.5701 - val_loss: 1.1720 - val_accuracy: 0.5557 - lr: 0.0010
```

```
Epoch 6/48
```

```
225/225 [=====] - 29s 128ms/step - loss: 1.0749 - accuracy:
0.5948 - val_loss: 1.2184 - val_accuracy: 0.5219 - lr: 0.0010
```

```
Epoch 7/48
```

```
225/225 [=====] - ETA: 0s - loss: 1.0297 - accuracy: 0.6094
Restoring model weights from the end of the best epoch: 4.
```

```
Epoch 7: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.
```

```
225/225 [=====] - 33s 147ms/step - loss: 1.0297 - accuracy:
0.6094 - val_loss: 1.2083 - val_accuracy: 0.5433 - lr: 0.0010
```

```
Epoch 7: early stopping
```

```
In [8]: plt.figure(figsize=(20,10))
plt.subplot(1, 2, 1)
plt.suptitle('Optimizer : Adam', fontsize=10)
plt.ylabel('Loss', fontsize=16)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend(loc='upper right')

plt.subplot(1, 2, 2)
plt.ylabel('Accuracy', fontsize=16)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
```

```
plt.legend(loc='lower right')  
plt.show()
```

Optimizer : Adam

